



GPU的CUDA编程方法

Floating-Point Operations per Second for the CPU and GPU



Memory Bandwidth for the CPU and GPU



Why a GPU?

CPU


Optimized for
Serial Tasks



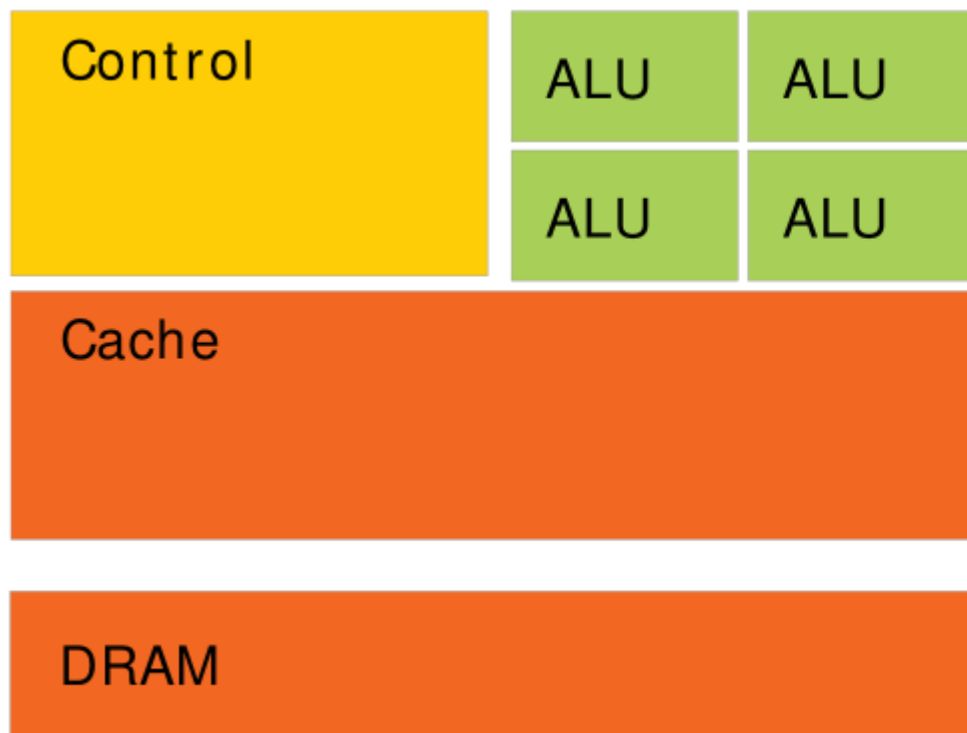
+

GPU

Optimized for Many
Parallel Tasks



The GPU Devotes More Transistors to Data Processing

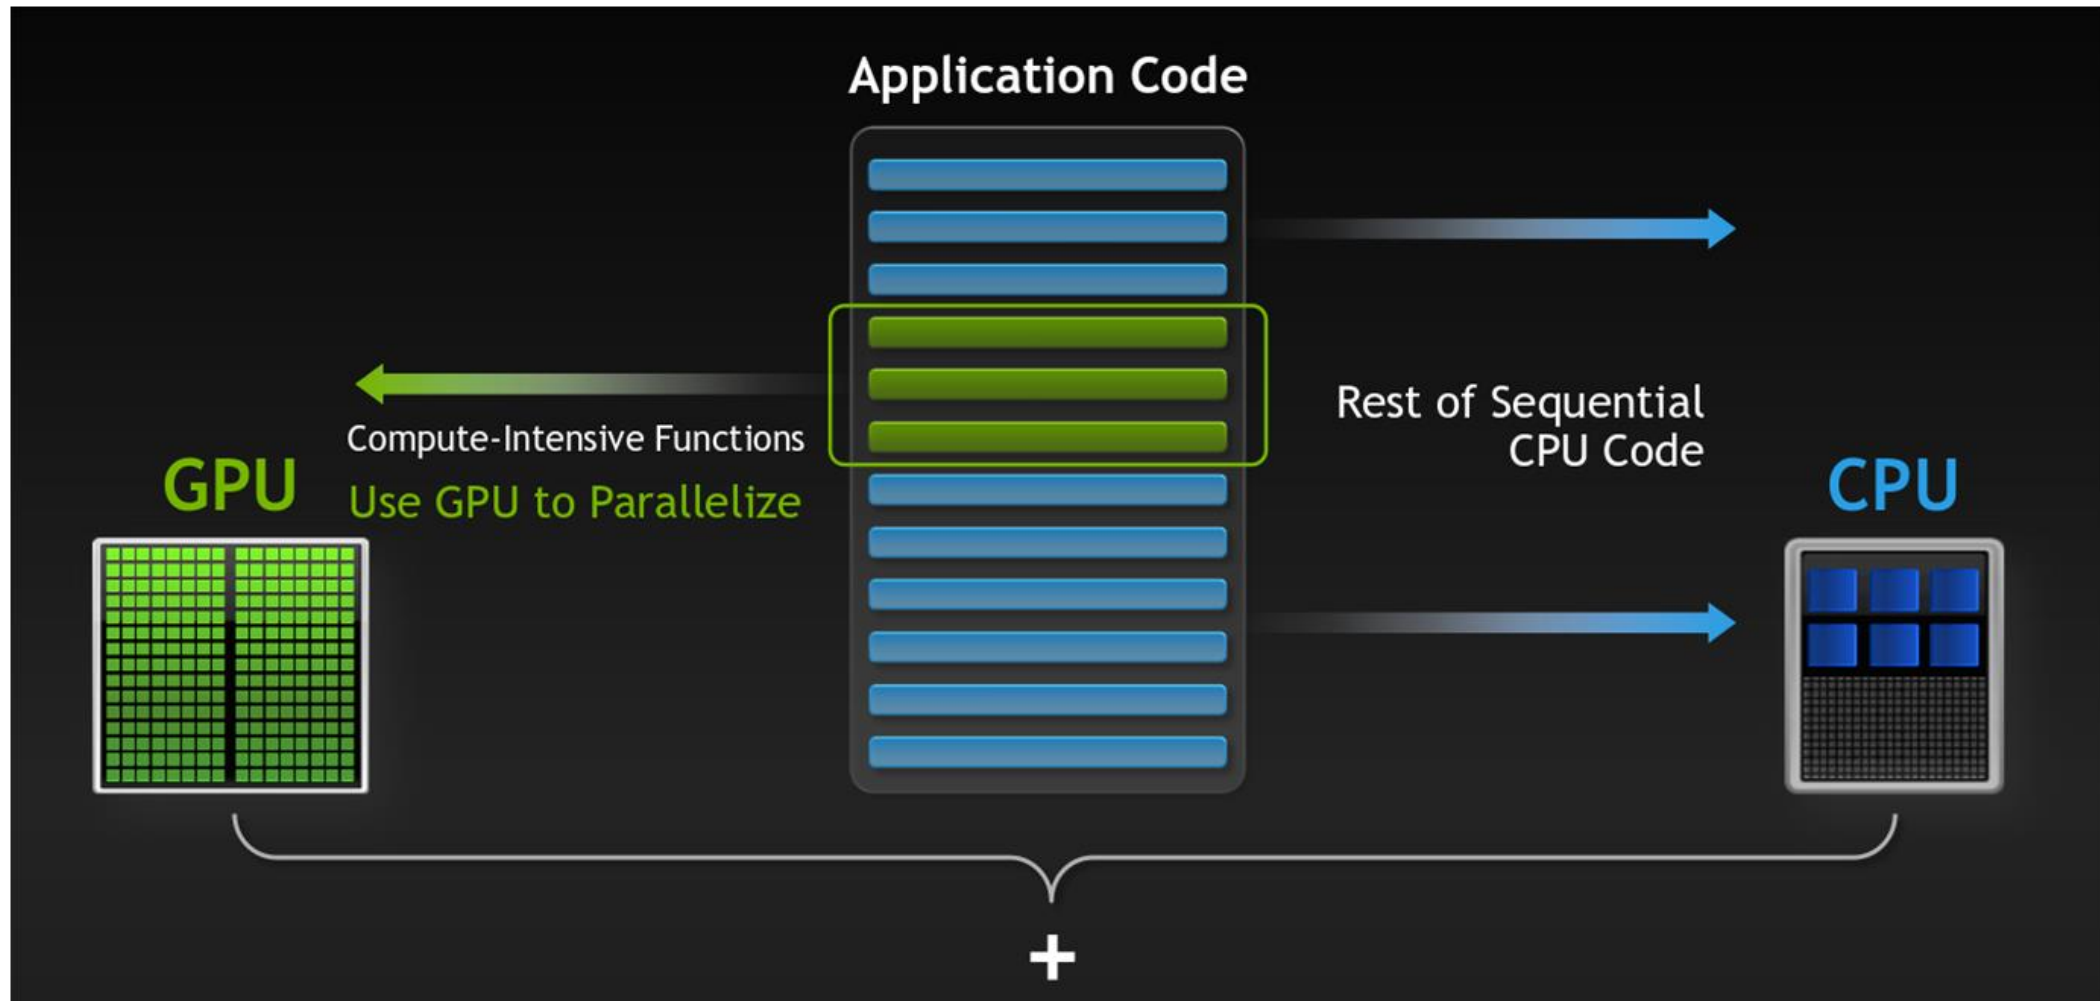


CPU



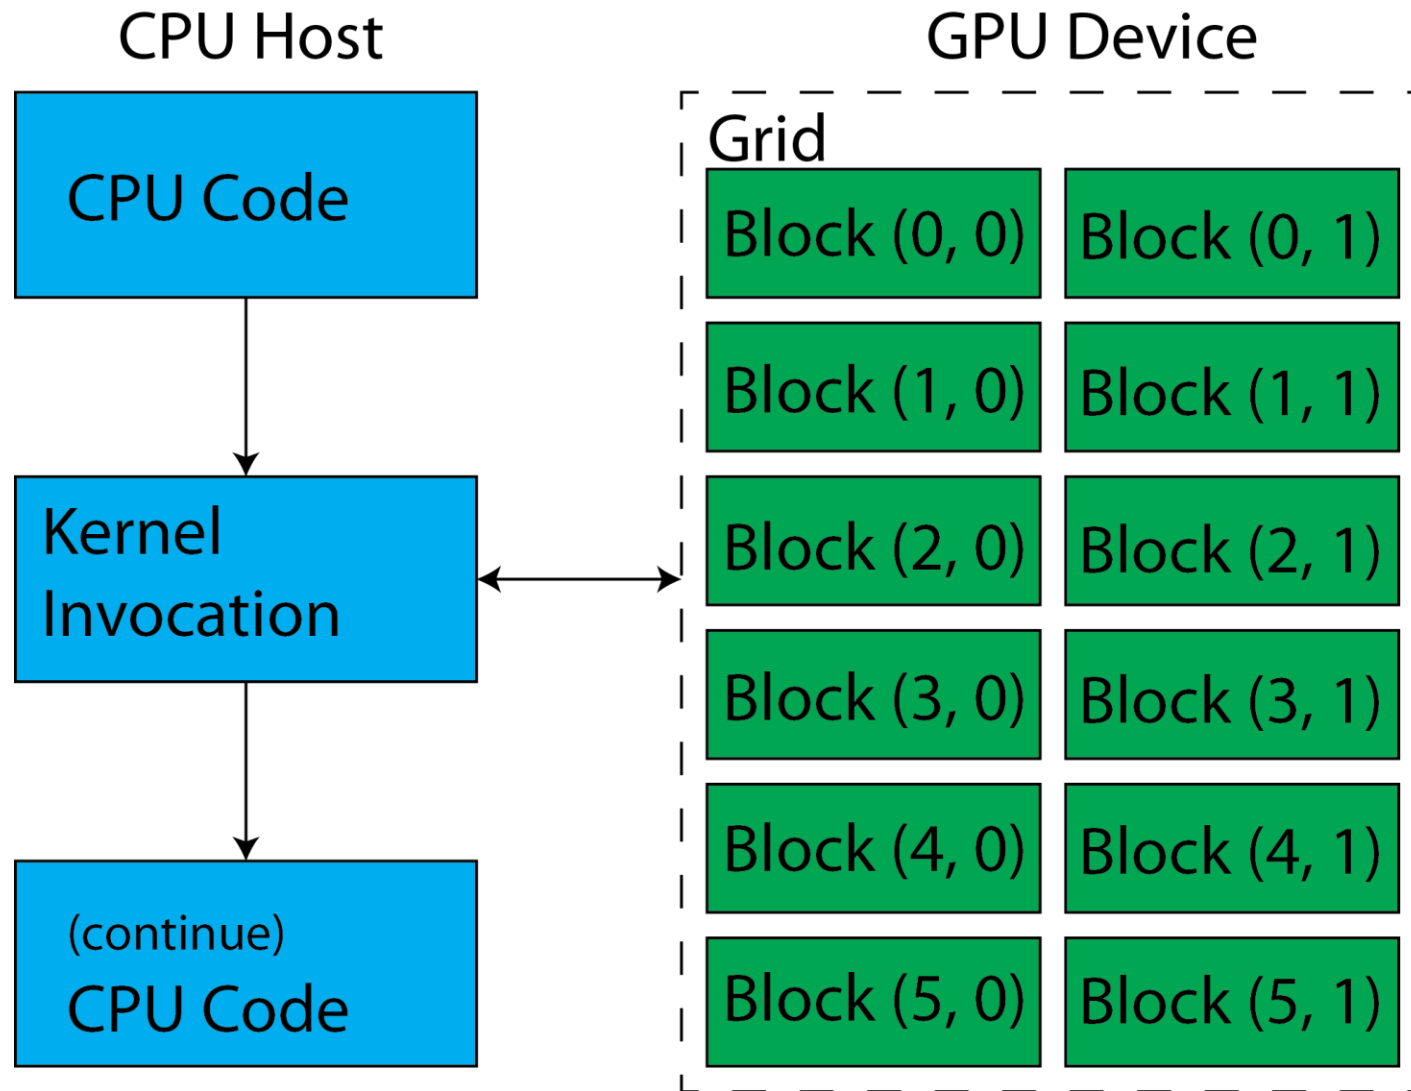
GPU

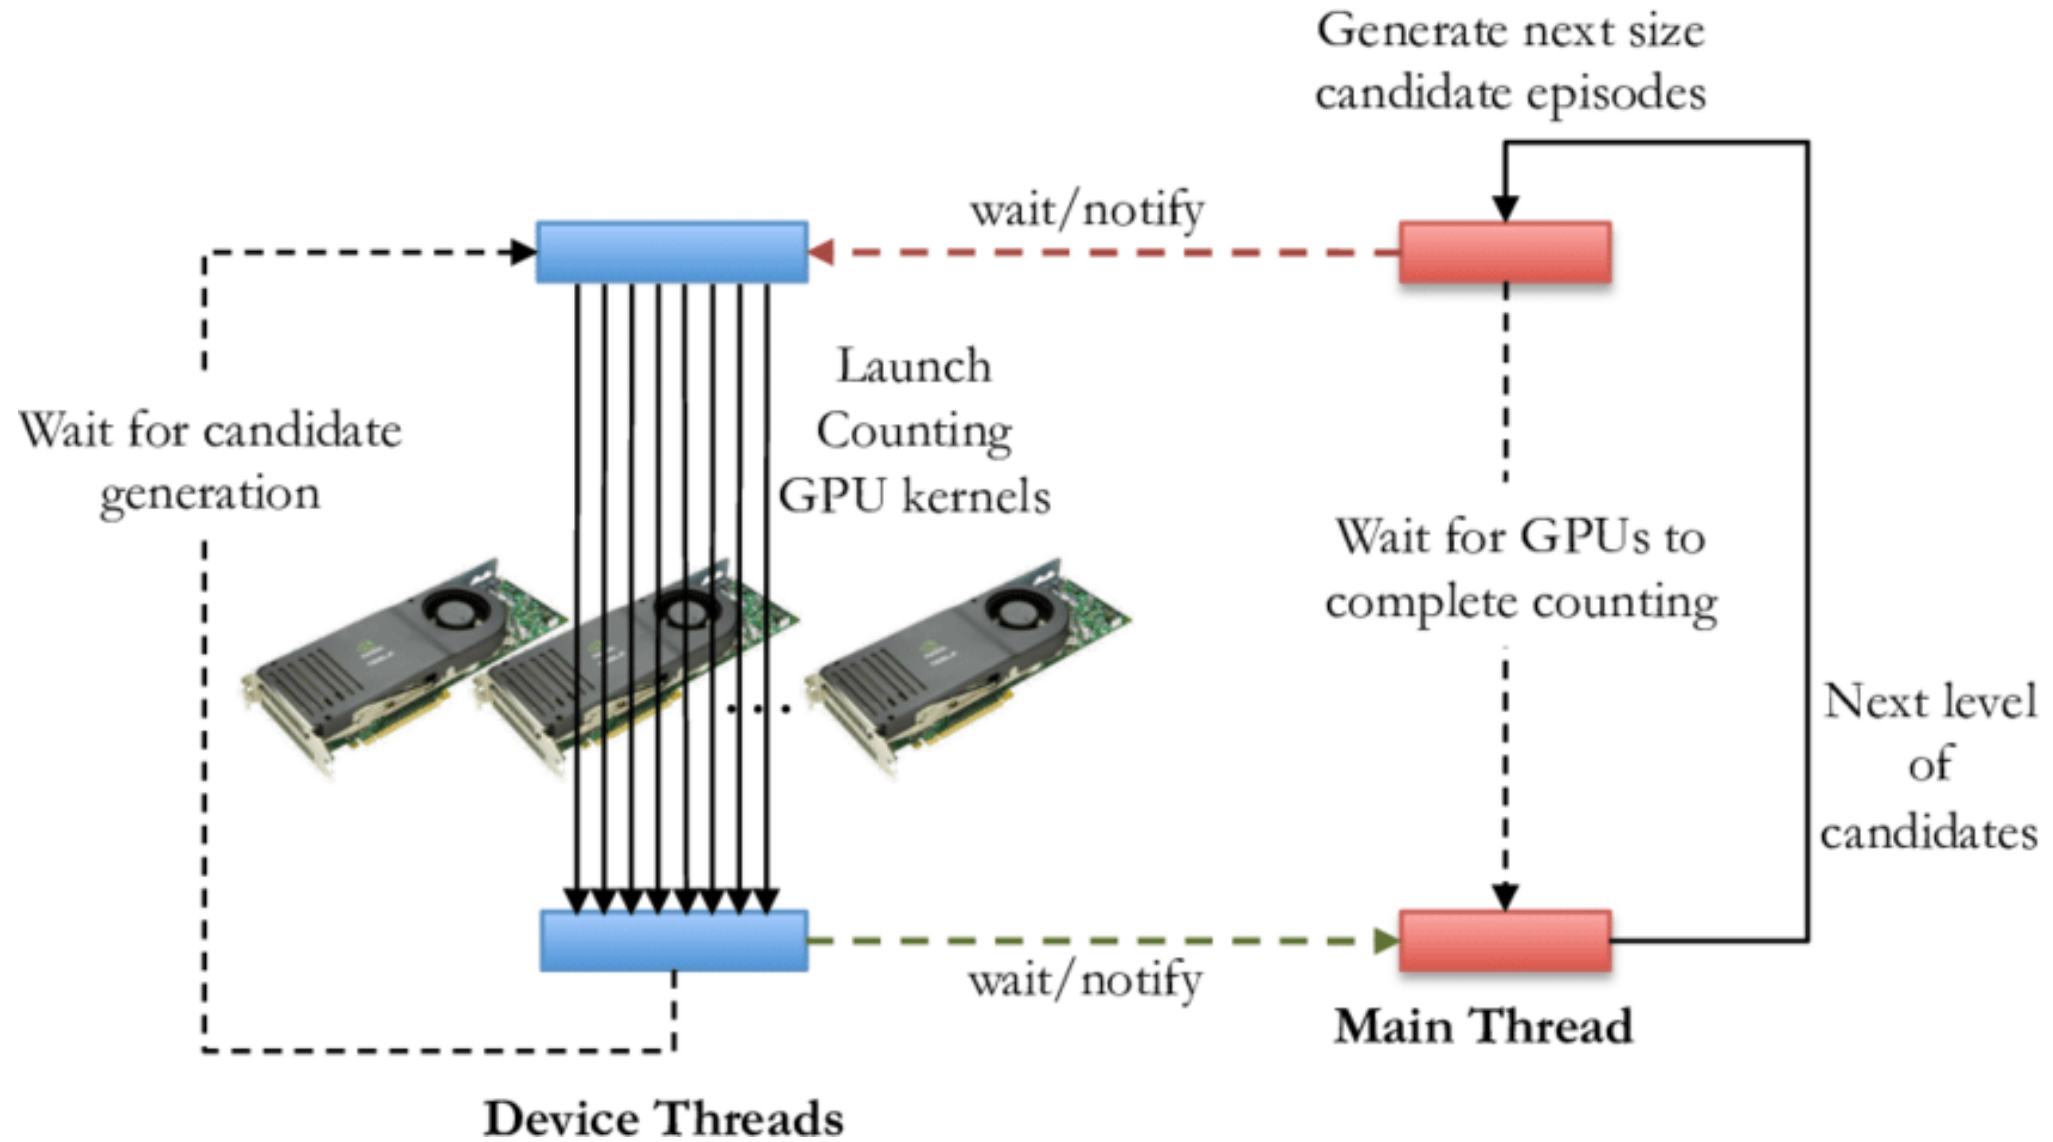
CUDA编程模型是一个异构模型，需要CPU和GPU协同工作

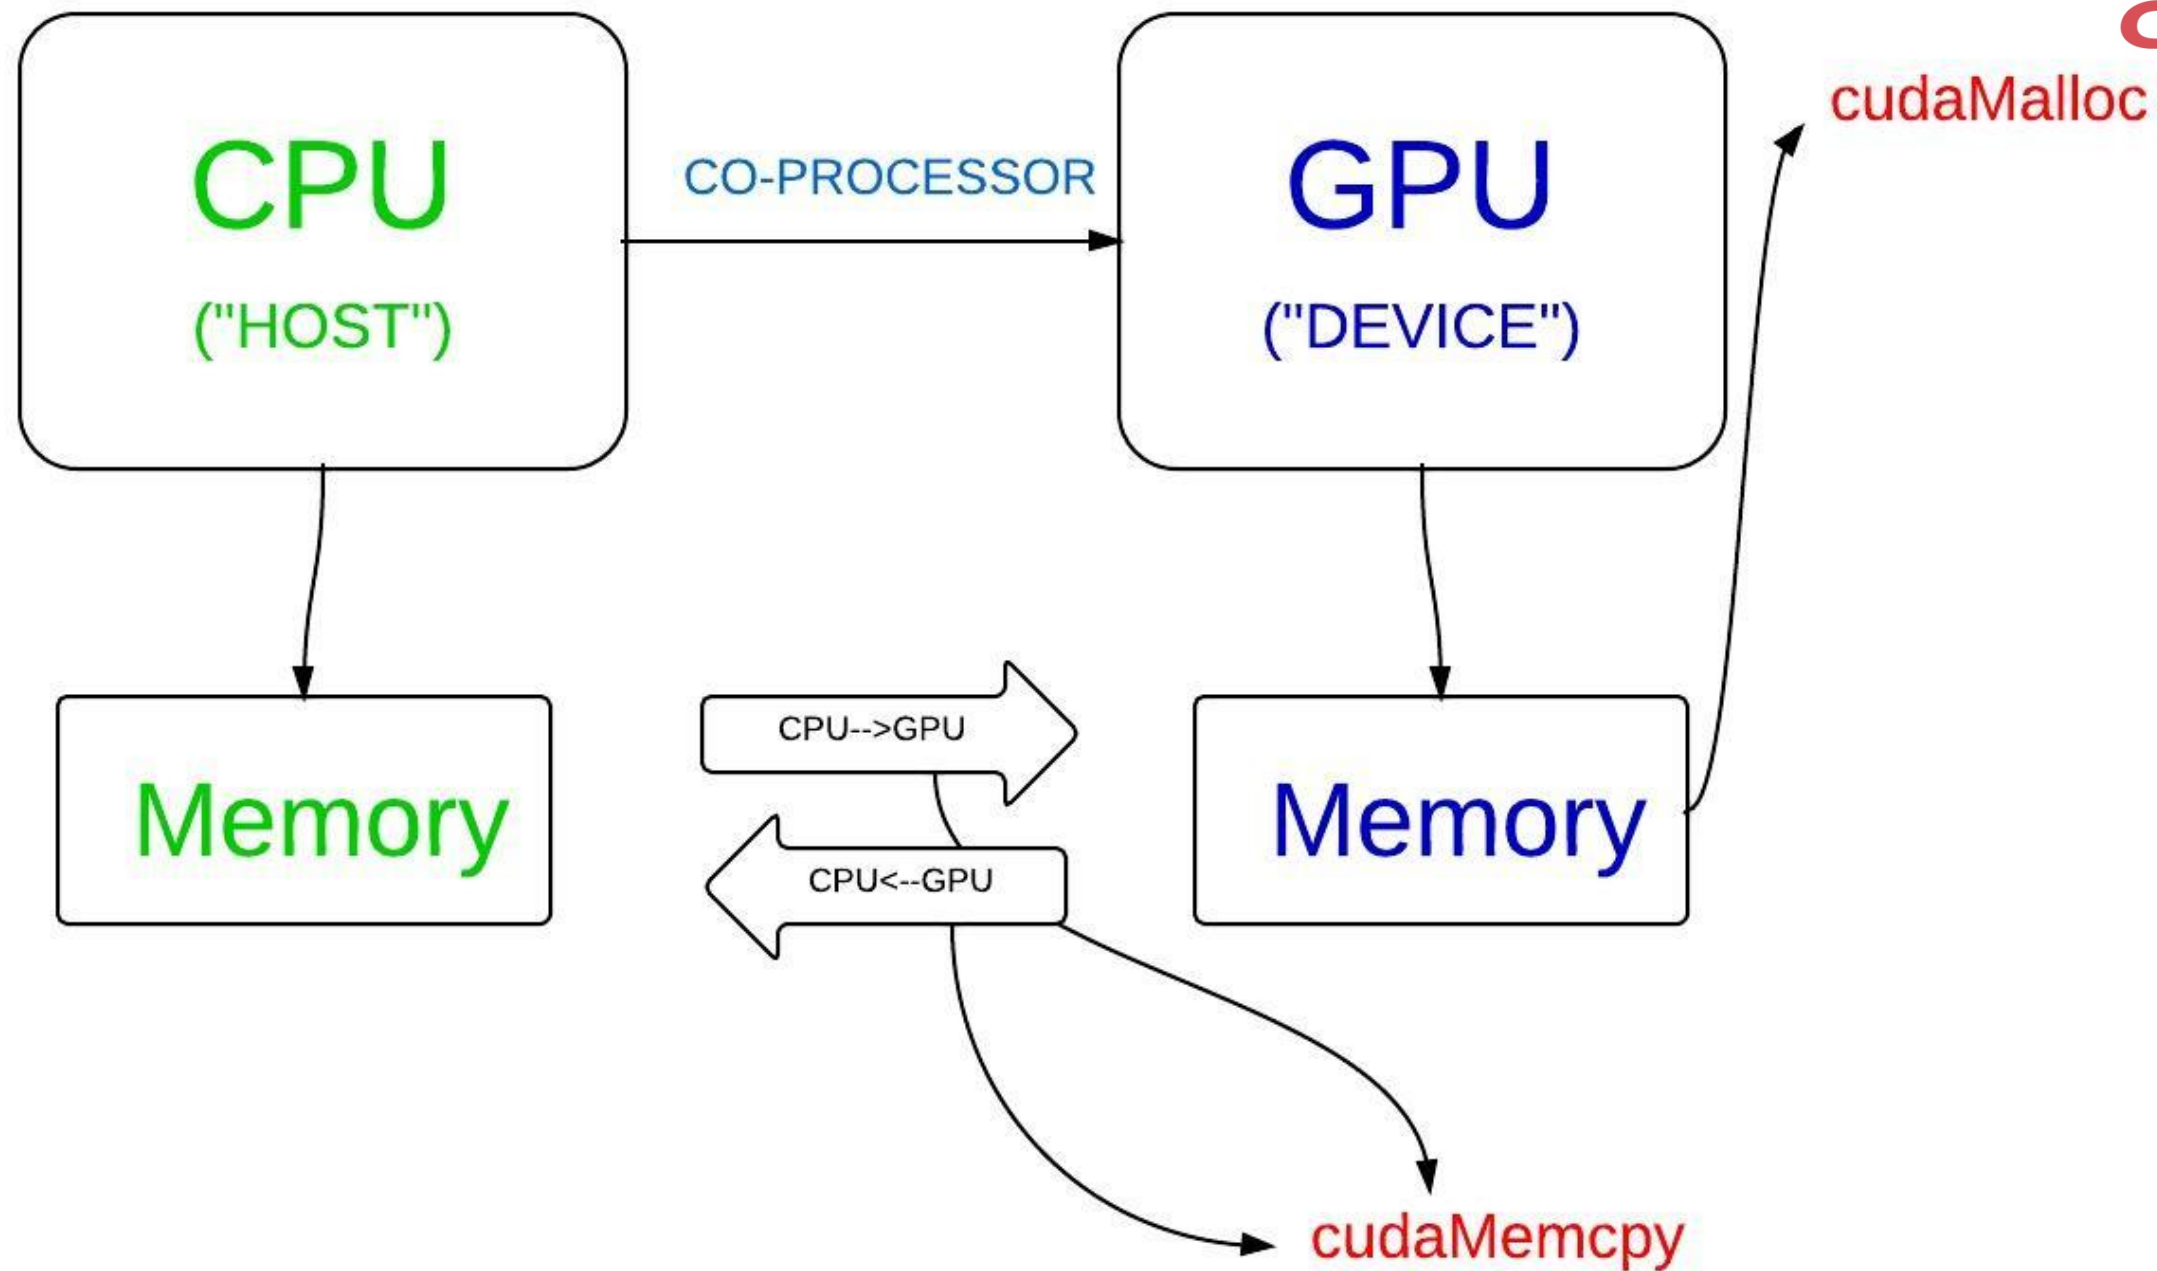


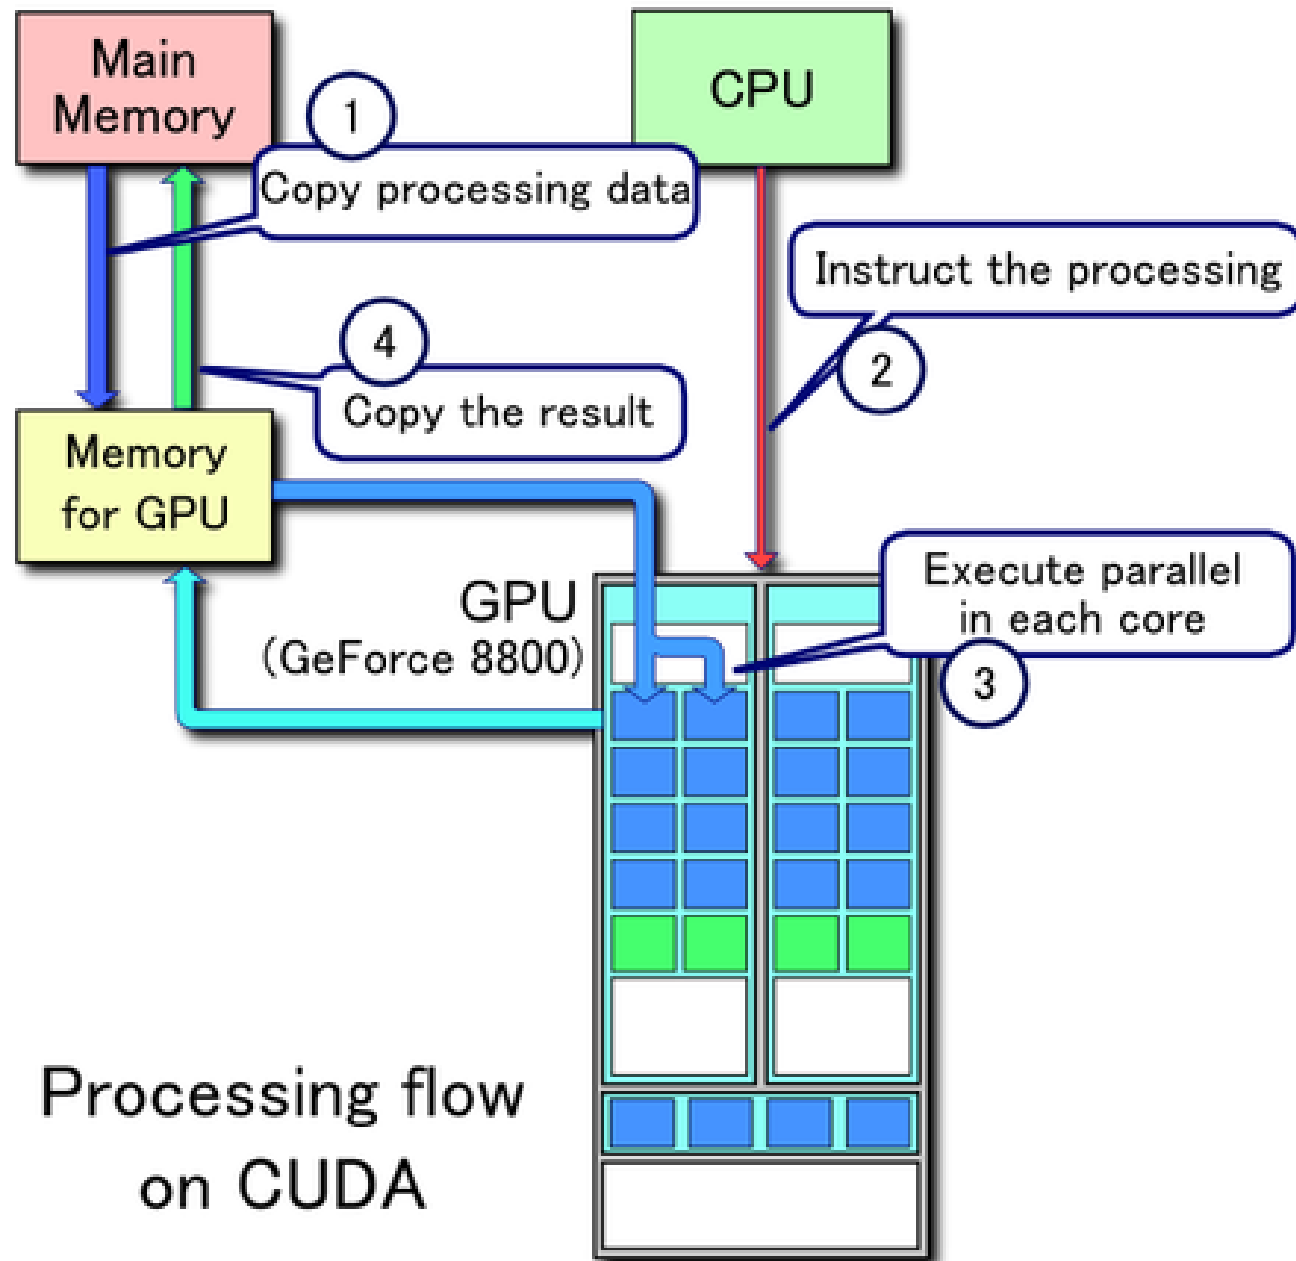
CUDA ® : A General-Purpose Parallel Computing Platform and Programming Model

GPU Computing Applications						
Libraries and Middleware						
cuDNN TensorRT	cuFFT, cuBLAS, cuRAND, cuSPARSE	CULA MAGMA	Thrust NPP	VSIPL, SVM, OpenCurrent	PhysX, OptiX, iRay	MATLAB Mathematica
Programming Languages						
C	C++	Fortran	Java, Python, Wrappers	DirectCompute	Directives (e.g., OpenACC)	
CUDA-enabled NVIDIA GPUs						
Turing Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier	GeForce 2000 Series		Quadro RTX Series	Tesla T Series	
Volta Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier				Tesla V Series	
Pascal Architecture (Compute capabilities 6.x)	Tegra X2	GeForce 1000 Series		Quadro P Series	Tesla P Series	
Maxwell Architecture (Compute capabilities 5.x)	Tegra X1	GeForce 900 Series		Quadro M Series	Tesla M Series	
Kepler Architecture (Compute capabilities 3.x)	Tegra K1	GeForce 700 Series GeForce 600 Series		Quadro K Series	Tesla K Series	
	EMBEDDED	CONSUMER DESKTOP, LAPTOP		PROFESSIONAL WORKSTATION	DATA CENTER	

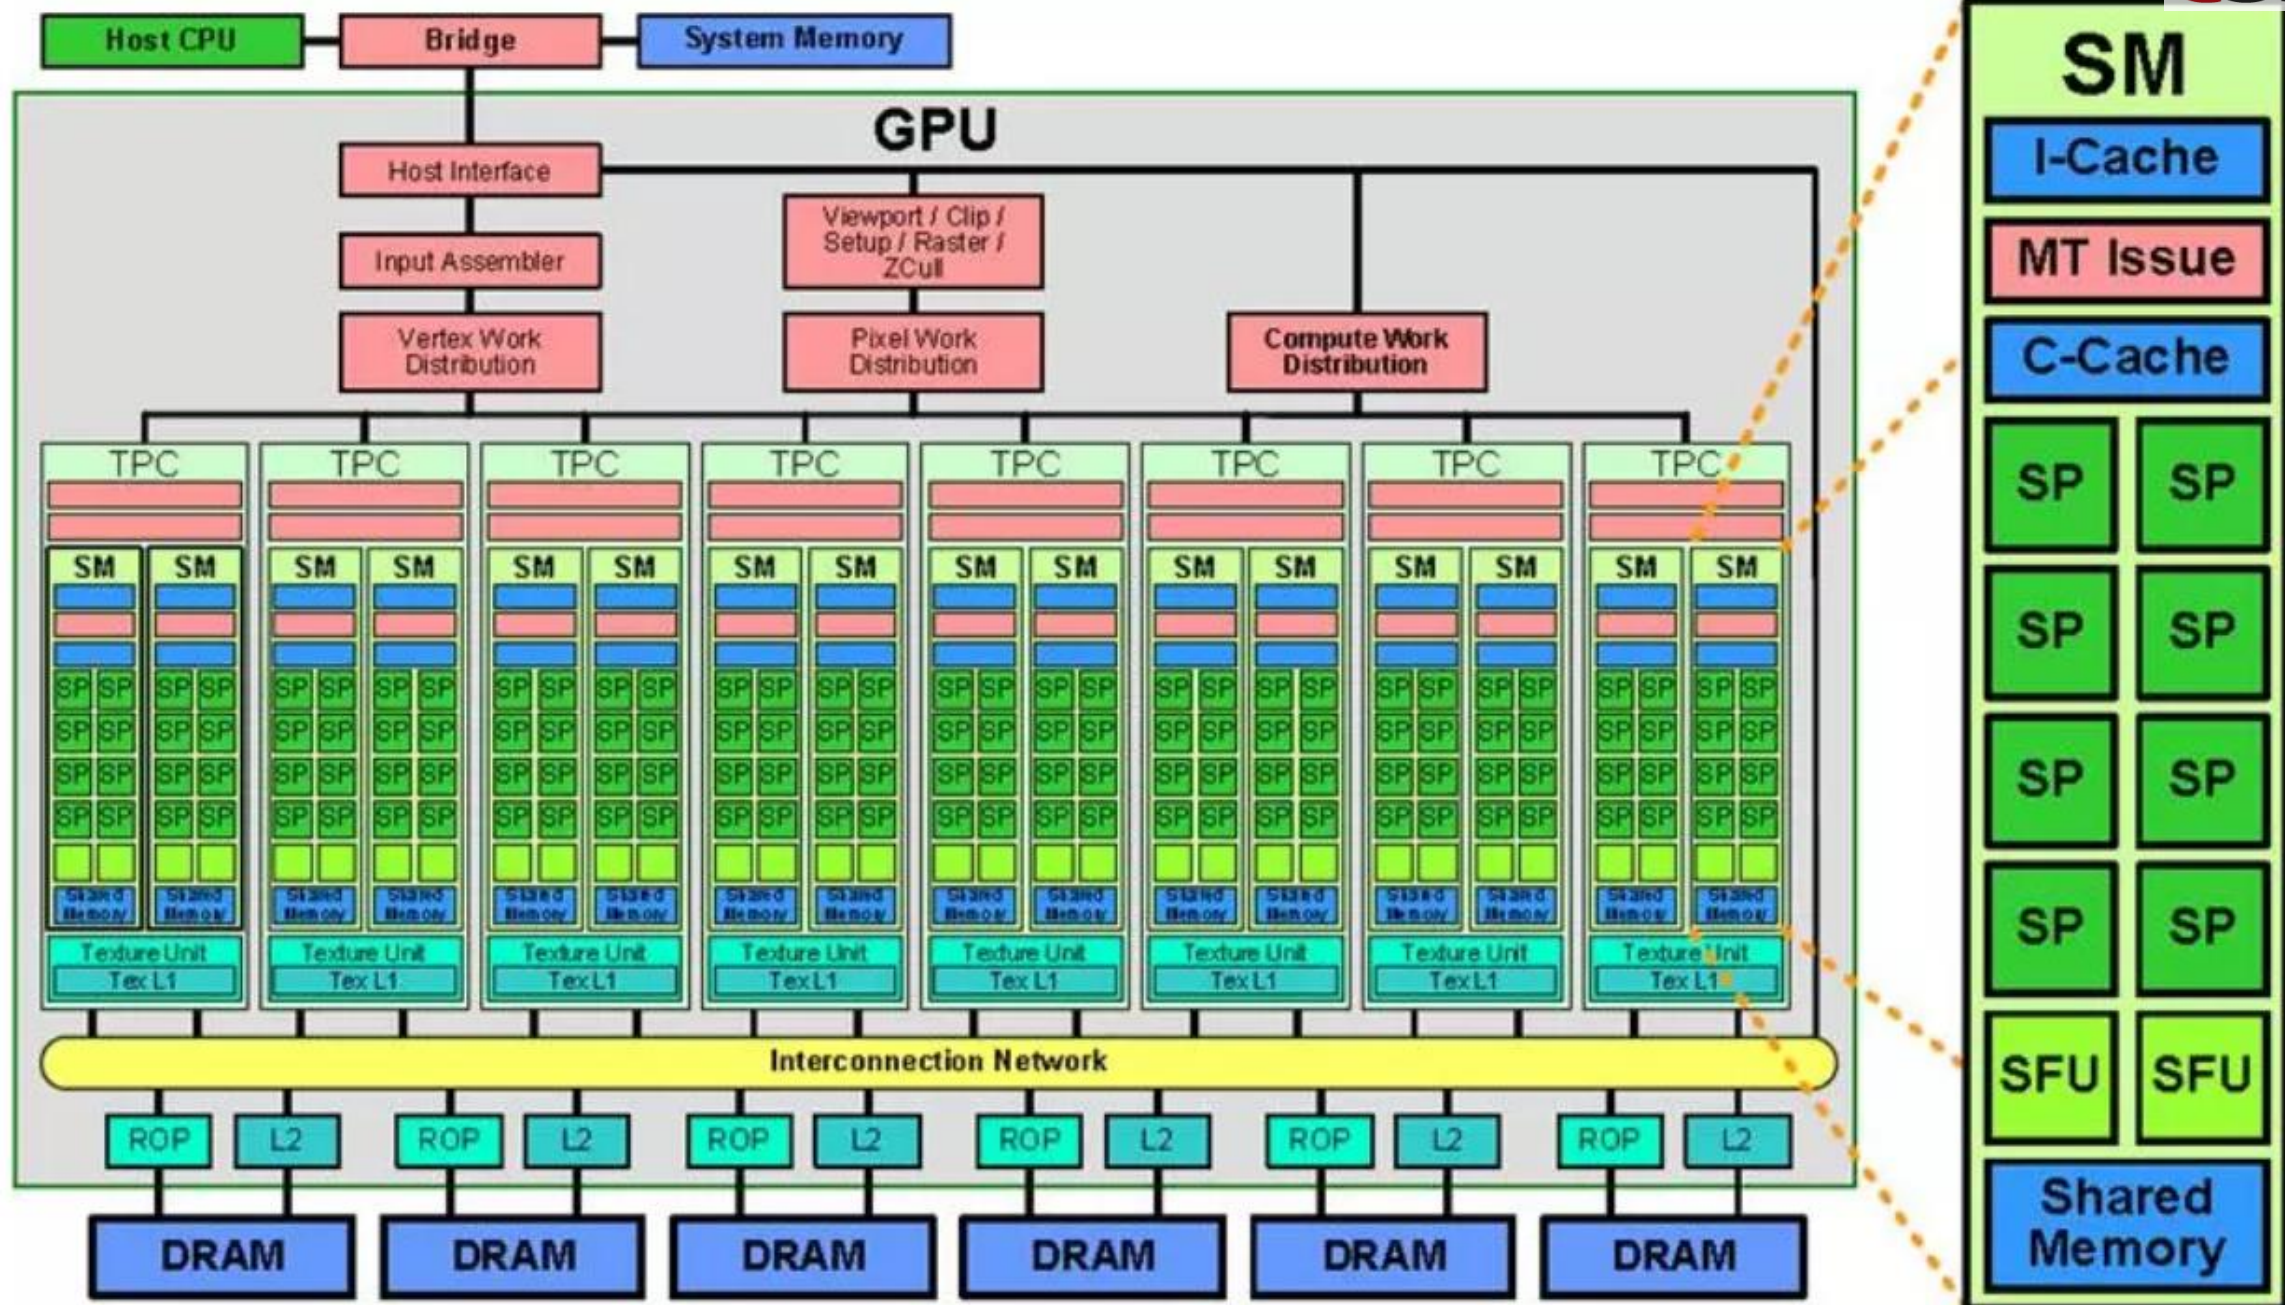








Processing flow
on CUDA

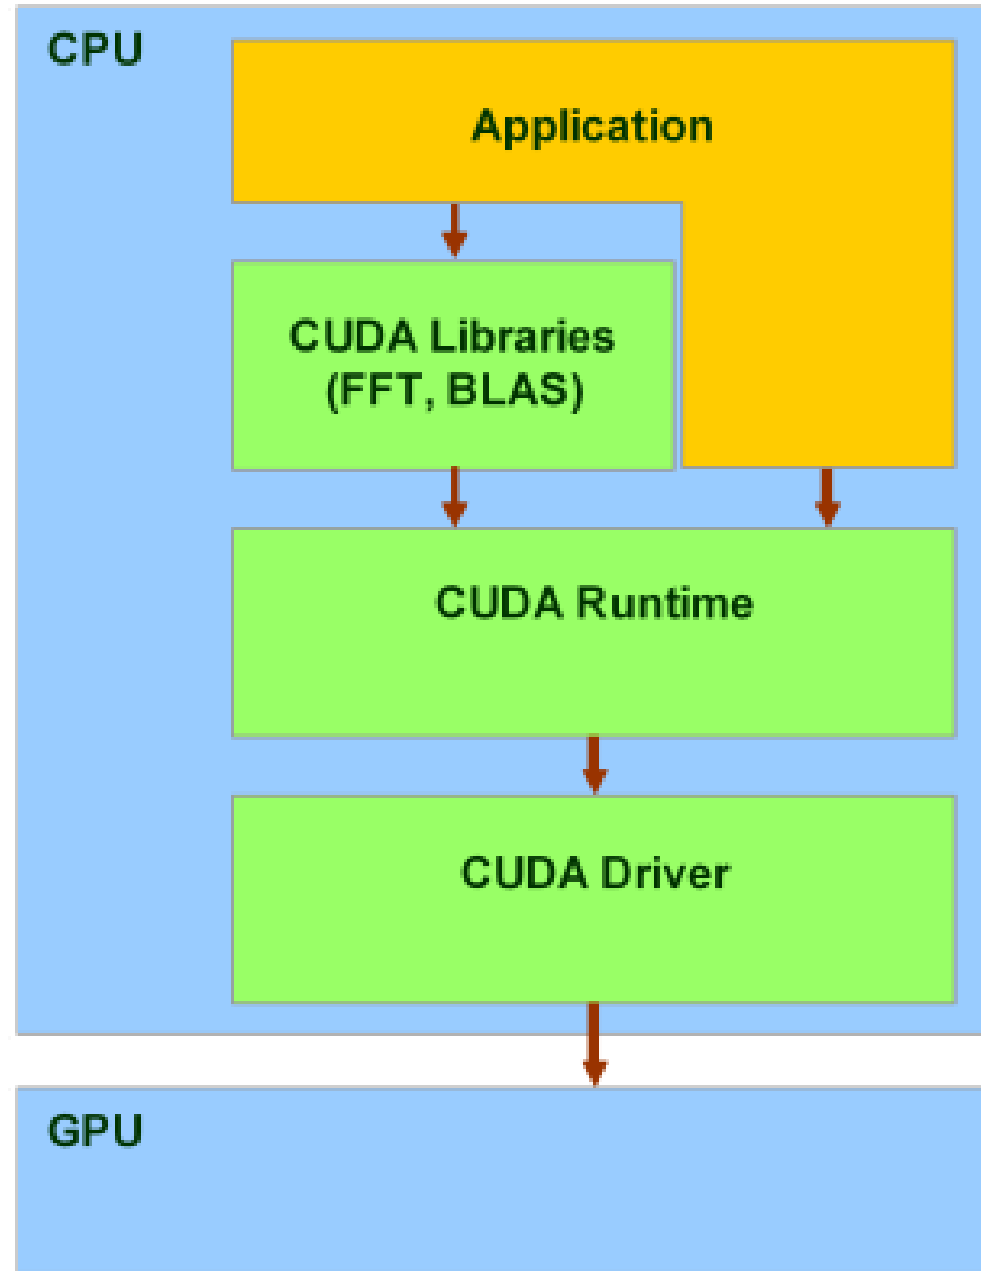


CPU	GPU	层次
算术逻辑和控制单元	流处理器(SM)	硬件
算术单元	批量处理器(SP)	硬件
进程	Block	软件
线程	thread	软件
调度单位	Warp	软件

SP最基本的处理单元, Streaming Processor, 也称为CUDA core。

SM是英文名是 Streaming Multiprocessor, 翻译过来就是流式多处理器。

SM采用的是SIMT (Single-Instruction, Multiple-Thread, 单指令多线程)架构, 基本的执行单元是warps, 一个warp包含32个线程。



CUDA Software Development

CUDA Optimized Libraries:
math.h, FFT, BLAS, ...

Integrated CPU + GPU
C Source Code

NVIDIA C Compiler

NVIDIA Assembly
for Computing (PTX)

CPU Host Code

CUDA
Driver

Profiler

Standard C Compiler

GPU

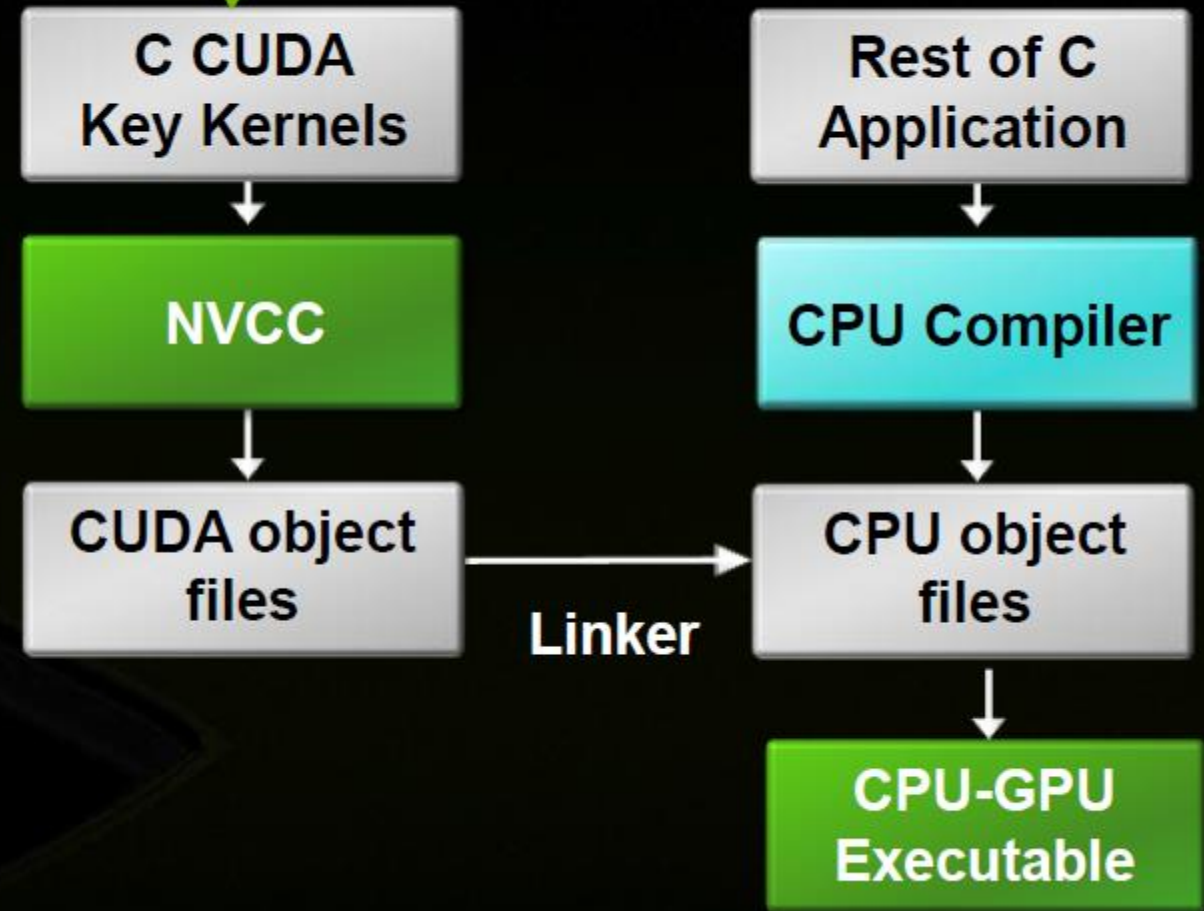
CPU

```
void serial_function(... ) {  
    ...  
}  
void other_function(int ... ) {  
    ...  
}
```

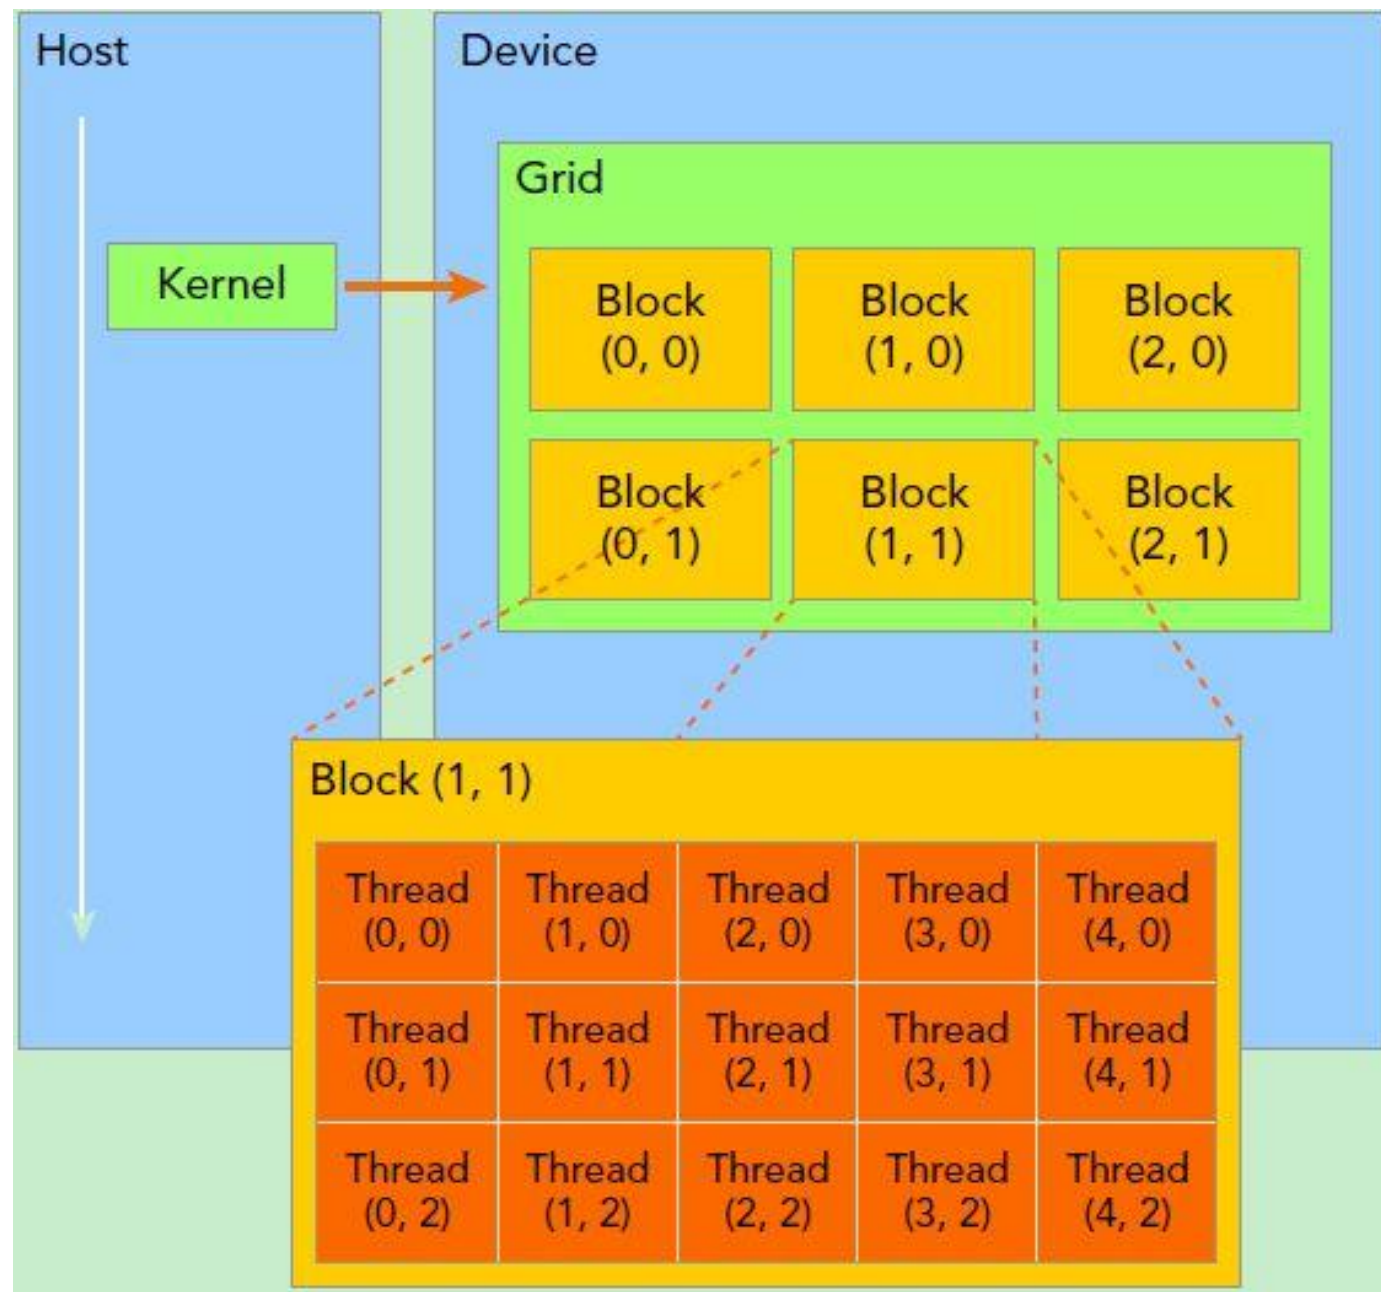
```
void saxpy_serial(float ... ) {  
    for (int i = 0; i < n; ++i)  
        y[i] = a*x[i] + y[i];  
}
```

```
void main( ) {  
    float x;  
    saxpy_serial(..);  
    ...  
}
```

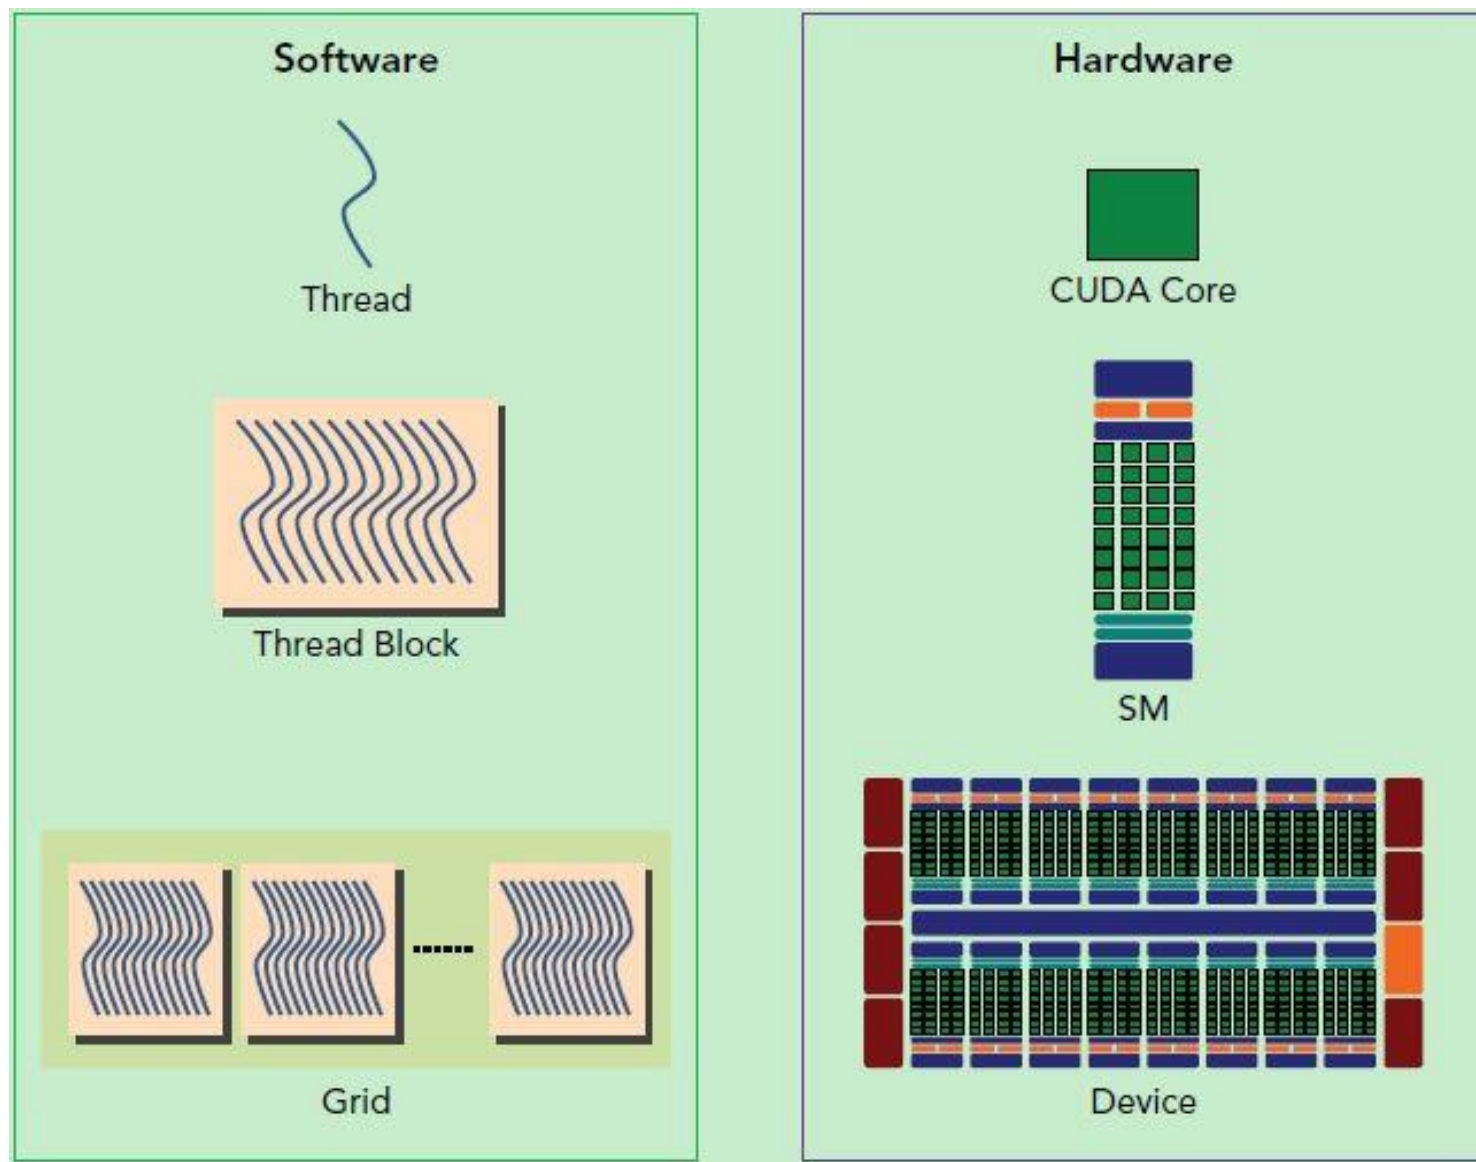
Write Parallel
CUDA code



Grid of Thread Blocks




Kernel上的两层线程组织结构 (2-dim)



- 每个thread由每个SP执行
- 每个thread block由SM执行

```
// Kernel definition
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
}

int main()
{
    ...
    // Kernel invocation with N threads
    VecAdd<<<1, N>>>(A, B, C);
    ...
}
```




指定kernel要执行的线程数量

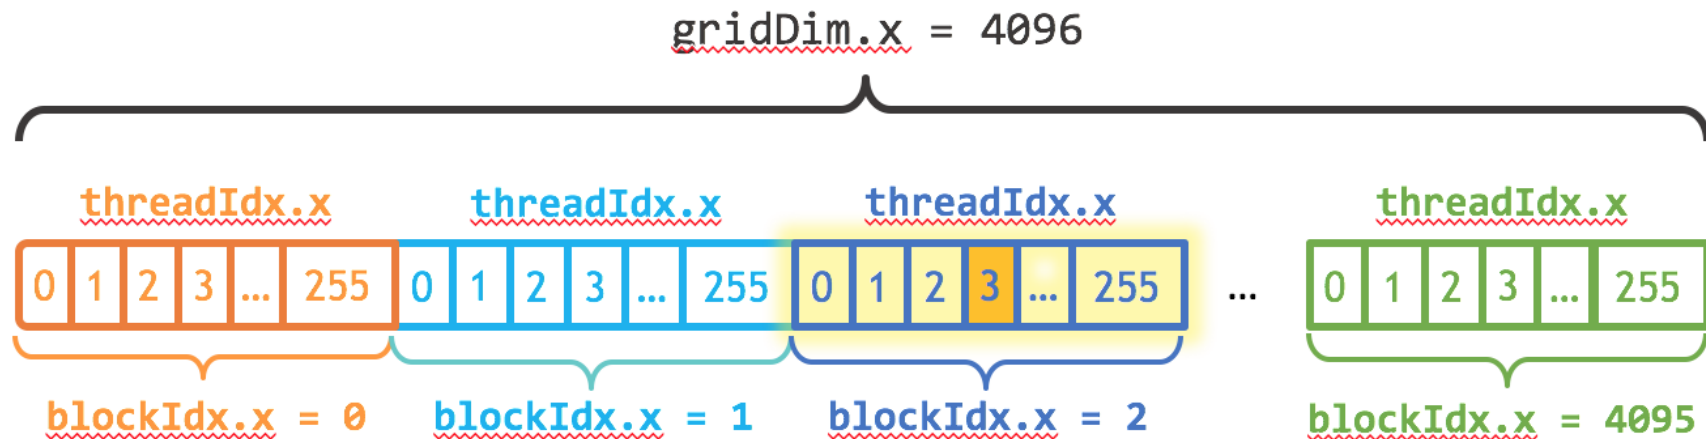
<<<1, size>>>, 表示分配了一个线程块 (Block), 每个线程块有分配了size个线程

"<<<>>>" 中的 参数并不是传递给设备代码的参数, 而是定义主机代码运行时如何启动设备代码。

CUDA C语言对C语言的扩展之一就是加入了一些函数前缀

	Executed on the:	Only callable from the:
__device__ float DeviceFunc()	device	device
__global__ void KernelFunc()	device	host
__host__ float HostFunc()	host	host

若向量大小为 $1 \ll 20$ ，而block大小为256，那么grid大小是4096，kernel的线程层级结构如下图所示：




$$\text{index} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$$


$$\text{index} = (2) * (256) + (3) = 515$$

- 一个线程需要两个内置的坐标变量 (blockIdx, threadIdx) 来唯一标识，它们都是dim3类型变量，其中 **blockIdx** 指明线程所在grid中的位置，而 **threadIdx** 指明线程所在block中的位置。
- 对于一个2-dim的block(Dx,Dy)，线程(x,y)的ID值为(x+y*Dx)，如果是3-dim的block(Dx,Dy,Dz)，线程(x,y,z)的ID值为(x+y*Dx+z*Dx*Dy)。
- 不是block越大越好，而要适当选择。目前GPU可以达到 **1024** threads/block

1D线程维度



2D线程维度



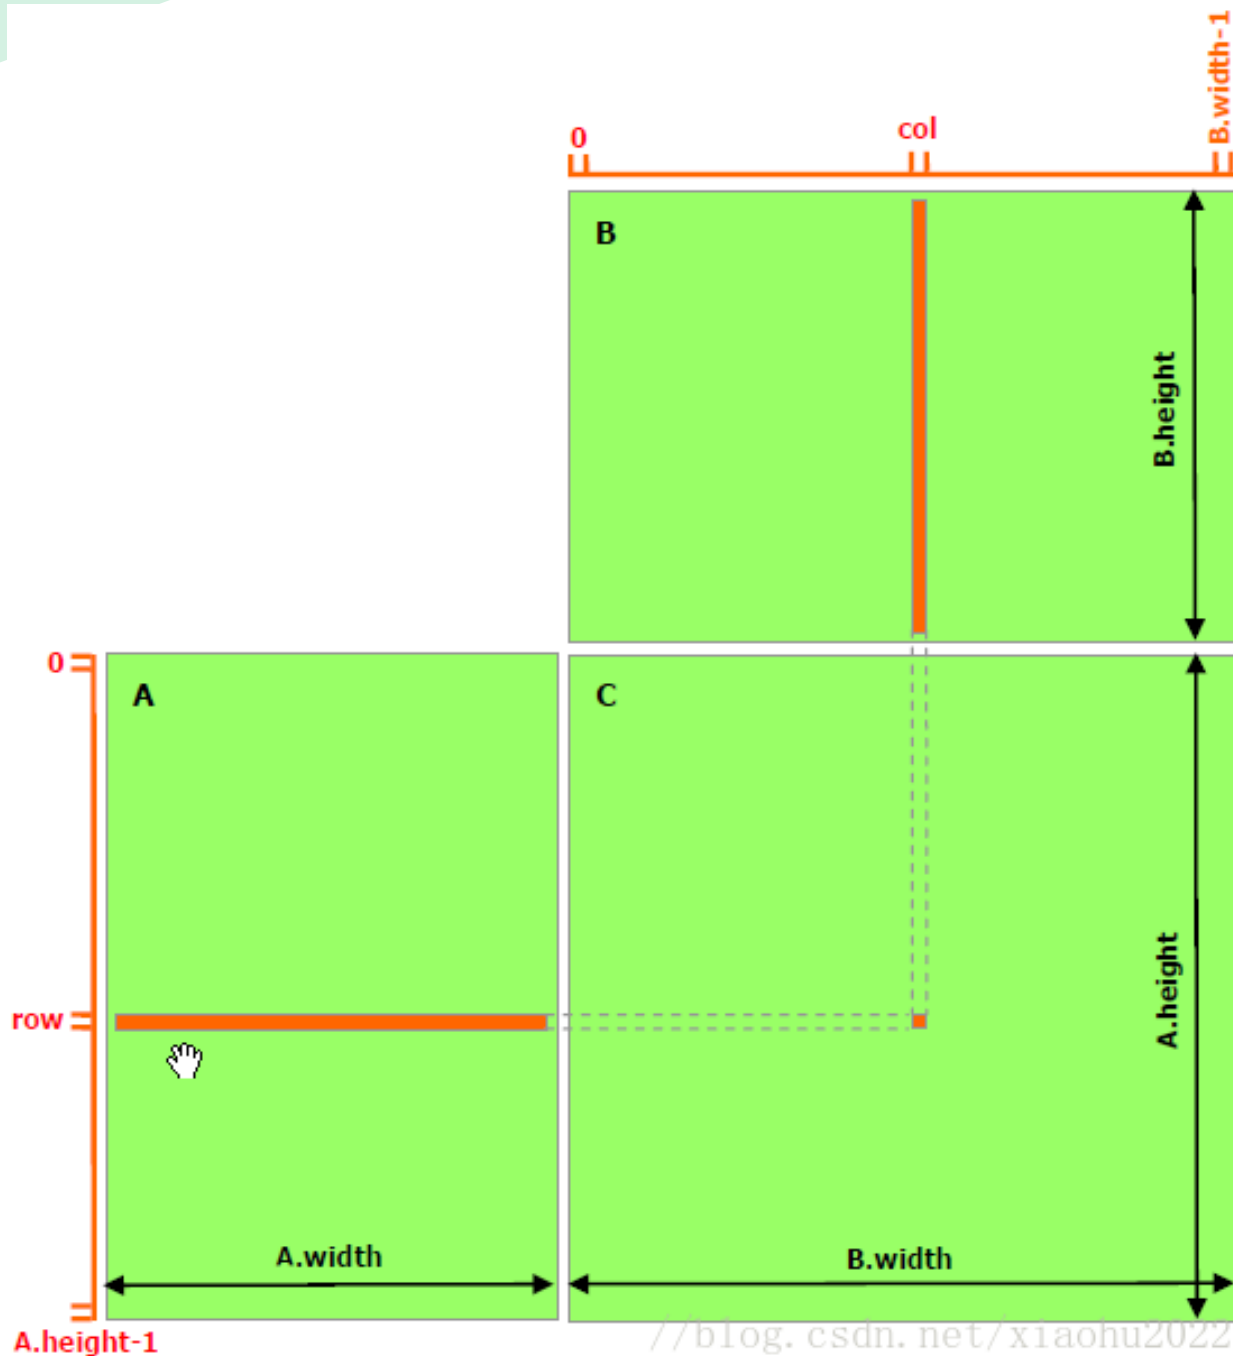
```
// Kernel definition
__global__ void MatAdd(float A[N][N], float B[N][N],
                      float C[N][N])
{
    int i = threadIdx.x;
    int j = threadIdx.y;
    C[i][j] = A[i][j] + B[i][j];
}

int main()
{
    ...
    // Kernel invocation with one block of N * N * 1 threads
    int numBlocks = 1;
    dim3 threadsPerBlock(N, N);
    MatAdd<<<numBlocks, threadsPerBlock>>>(A, B, C);
    ...
}
```


kernel的这种线程组织结构天然适合vector, matrix等运算, 如利用2-dim结构实现两个矩阵的加法, 每个线程负责处理每个位置的两个元素相加, 代码如下所示。线程块大小为(16, 16), 然后将 $N \times N$ 大小的矩阵均分为不同的线程块来执行加法运算。


```
// Kernel定义
__global__ void MatAdd(float A[N][N], float B[N][N], float C[N][N])
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;
    if (i < N && j < N)
        C[i][j] = A[i][j] + B[i][j];
}

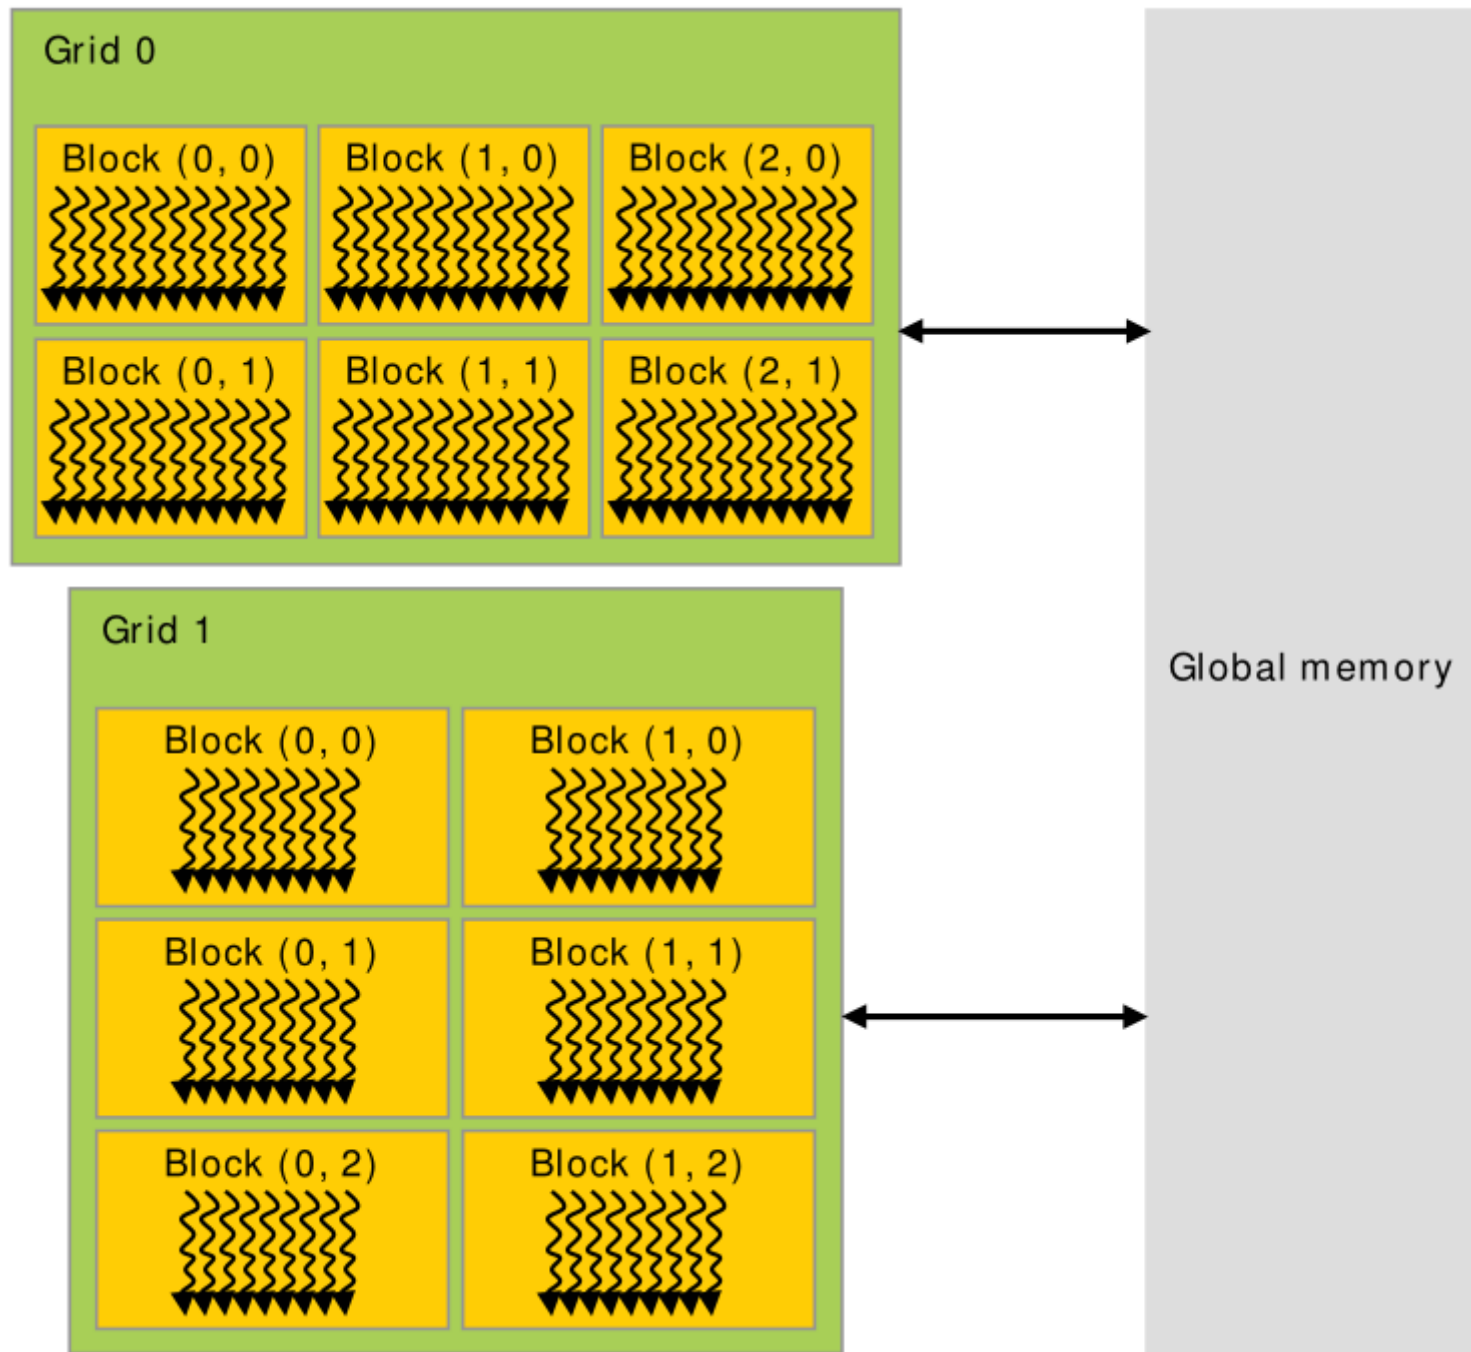
int main()
{
    ...
    // Kernel 线程配置
    dim3 threadsPerBlock(16, 16);
    dim3 numBlocks(N / threadsPerBlock.x, N / threadsPerBlock.y);
    // kernel调用
    MatAdd<<<numBlocks, threadsPerBlock>>>(A, B, C);
    ...
}
```



这里矩阵大小为 1024×1024 ，设计的线程的block大小为 $(32, 32)$ ，那么grid大小为 $(32, 32)$ 。

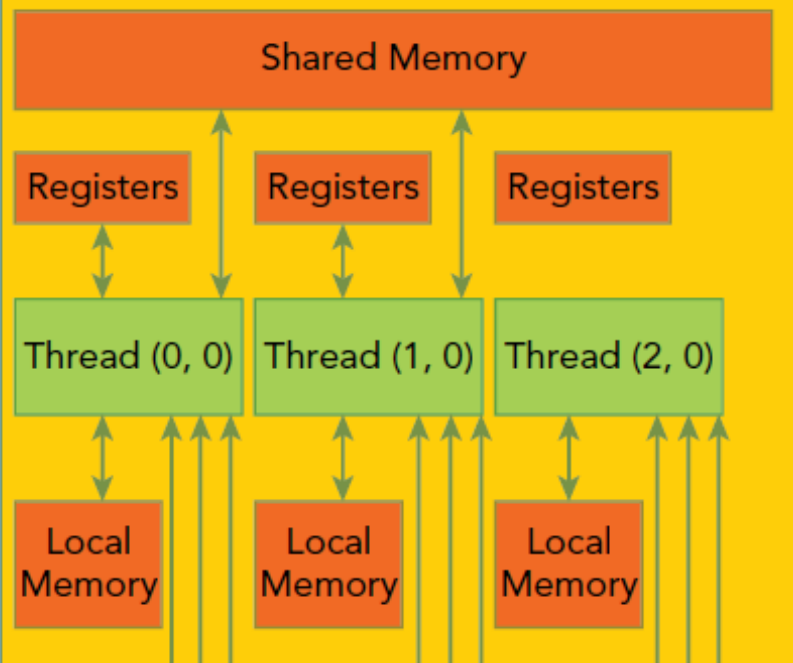
CUDA的内存模型





(Device) Grid

Block (0, 0)



每个线程有自己的私有本地内存 (Local Memory)，而每个线程块有包含共享内存 (Shared Memory)，可以被线程块中所有线程共享，其生命周期与线程块一致。此外，所有的线程都可以访问全局内存 (Global Memory)。还可以访问一些只读内存块：常量内存 (Constant Memory) 和纹理内存 (Texture Memory)。

Host

Global
MemoryConstant
MemoryTexture
Memory

C Program
Sequential
Execution

Serial code

Parallel kernel

Kernel0<<<>>>()

Serial code

Parallel kernel

Kernel1<<<>>>()

Host



Device

Grid 0

Block (0, 0)



Block (1, 0)



Block (2, 0)



Block (0, 1)



Block (1, 1)



Block (2, 1)



Host



Device

Grid 1

Block (0, 0)



Block (1, 0)



Block (0, 1)




Block (1, 1)



Block (0, 2)




Block (1, 2)



Threads的同步

Barrier – point in the program where threads stop and wait.

When all threads have reached the barrier, they can proceed.



```
Int idx = threadIdx.x;
__shared__ int array[128];
array[idx] = threadIdx.x;
__syncthreads();
If (idx < 127) {
    int temp = array[idx+1];
    __syncthreads();
    array[idx] = temp;
    __syncthreads();
}
```


CUDA synchronization

CUDA provides a synchronization **barrier** routine for those threads within each block
`__syncthreads()`

This routine would be used **within a kernel**.

Threads would wait at this point until all threads in the block have reached it and they are all released.

Note: **only synchronizes with other threads in block**.

七个步骤

`cudaSetDevice(0);` //获取设备; 只有一个GPU时或默认使用0号GPU时可以省略

`cudaMalloc((void**) &d_a, sizeof(float)*n);` //分配显存

`cudaMemcpy(d_a, a, sizeof(float)*n, cudaMemcpyHostToDevice);` //数据传输 host to device

`gpu_kernel<<<blocks, threads>>>(**);` //kernel函数

`cudaMemcpy(a, d_a, sizeof(float)*n, cudaMemcpyDeviceToHost);` //数据传输 device to host

`cudaFree(d_a);` //释放显存空间

`cudaDeviceReset();` //重置设备; 可以省略

GPU的axpy的实现

blas_kernel.cu

```
__global__ void axpy_kernel(int N, float ALPHA, float *X, int OFFX, int INCX, float *Y, int OFFY, int INCY)
{
    int i = (blockIdx.x + blockIdx.y*gridDim.x) * blockDim.x + threadIdx.x;
    if(i < N) Y[OFFY+i*INCY] += ALPHA*X[OFFX+i*INCX];
}
```

```
extern "C" void axpy_ongpu(int N, float ALPHA, float * X, int INCX, float * Y, int INCY)
{
    axpy_ongpu_offset(N, ALPHA, X, 0, INCX, Y, 0, INCY);
}
```

```
extern "C" void axpy_ongpu_offset(int N, float ALPHA, float * X, int OFFX, int INCX, float * Y, int OFFY, int INCY)
{
    axpy_kernel<<<cuda_gridsize(N), BLOCK, 0, get_cuda_stream()>>>(N, ALPHA, X, OFFX, INCX, Y, OFFY, INCY);
    CHECK_CUDA(cudaPeekAtLastError());
}
```